

A Paradoxical, Self-Referential System of Logics

Four non-classical logics as resolutions of one universal diagonal,
with four machine-checked components

Vitaly Reznik

2026 — Version 1.1.0

Lean 4 witnesses: `github.com/inventor1975/ZTL` (`lean/ZTL.lean`, `lean/QuantumWitness.lean`,
`lean/Contextuality.lean`, `lean/JunctionWitness.lean`)

Verify from zero: §6 — four files, four commands, no dependencies

Paradox and Truth — the progenitors of Logic!

Abstract

We propose a reading of four non-classical logics — classical, intuitionistic / type-theoretic, quantum (orthomodular), and Zero-Trust Logic (ZTL) — as four disciplined responses to a single self-referential construction: the universal diagonal underlying the liar, Russell, Cantor, Gödel and Tarski (Lawvere 1969; Yanofsky 2003). Each retains all classical structural principles *save one*, and each thereby houses the diagonal differently. Two features are invariant across all four: non-contradiction with its associated reductio (a retained *principle*), and the diagonal itself (a retained *obstruction*). We isolate the parts of this picture that are formal rather than expository — the complementary sacrifice made by ZTL and by orthomodular quantum logic, and the combinatorial core of quantum contextuality (Mermin–Peres, GHZ) and the junction theorem of composition — and record them as four machine-checked components in Lean 4 on the *empty axiom list* (no classical choice, no quotient, no propositional extensionality). A closing section, explicitly marked as a reading, draws the consequence of the contextuality witness: the quantum “corner” pays its laws only in the DESCRIPTIONS; at the level of states the diagonal’s premise itself fails — the self-negating fixed point exists physically — so the paradox, for the world, dissolves rather than being tamed. We are explicit throughout about what is established (all four logics; the Lawvere–Yanofsky unification), what is the author’s (ZTL; the arrangement; the two witnesses), and what is analogy carrying no proof (the “cycle” itself). This is a synthesis and a reading, not a theorem, a merger, a new foundation, or a new field.

Keywords. self-reference; Lawvere fixed point; diagonal argument; quantum logic; orthomodular lattice; three-valued logic; paracomplete logic; Zero-Trust Logic; universal logic; Lean 4; machine-checked proof.

MSC 2020. 03B50; 03B60; 03G12; 18C50; 03B35.

1 Scope

The claim of this note is deliberately narrow. It is *not* that a new logic or a new foundation has been discovered, nor that the four logics below are amalgamated into one calculus — they are provably incompatible in the precise sense of §5. The unification of the classical antinomies under a single fixed-point schema is due to Lawvere [8] and, in the concrete set-theoretic form used here, to Yanofsky [14]; quantum logic is Birkhoff–von Neumann [1]; the intuitionistic corner is

standard; the “general theory of logical systems” already exists as a field (universal logic, abstract model theory [16]). The single component original to the author is **ZTL** [15], which appears here only as one corner. What this note adds is (i) the *arrangement* of the four as four resolutions of one diagonal, (ii) three machine-checked components (§6) — the two complementary poles and the combinatorial core of contextuality — and (iii) a closing reading (§7), marked as such. Everything expository is marked as such; everything proved is in §6 and reproducible from zero.

2 The universal diagonal

The classical antinomies are one construction in several presentations. Lawvere’s fixed-point theorem makes this exact.

Lemma 2.1 (Lawvere [8]). *In a cartesian closed category, if there is a point-surjective morphism $\varphi : A \rightarrow Y^A$, then every endomorphism $g : Y \rightarrow Y$ has a fixed point $y : 1 \rightarrow Y$ with $g \circ y = y$.*

Contrapositively: a *fixed-point-free* endomorphism of Y obstructs any point-surjection $A \rightarrow Y^A$. Instantiating Y and g :

- $Y = 2$, $g = \neg$ (no fixed point): no surjection $A \rightarrow 2^A$ — **Cantor**;
- the same schema, read in an internal truth object with a fixed-point-free negation, yields the **liar** and **Russell**;
- read in a provability / representation setting it yields **Gödel** and **Tarski**’s undefinability (Yanofsky [14], §§3–5).

The common core is *self-application against a fixed-point-free map*: an object that names its own maps, composed with a g that has no fixed point. Following Yanofsky, we call this schema **the diagonal**. It is *pre-classical*: it is not built by classical logic but is a construction that any logic expressive enough to talk about its own sentences must confront. A logic that ignores it does not escape it; it trivializes on it.

3 Four resolutions

A logic admitting self-reference must do something disciplined with the diagonal or collapse. There are (at least) four such disciplines in current use, and each is *classical logic minus exactly one structural principle*:

Logic	Principle dropped	Resolution of the diagonal
Classical	none	explodes on the raw liar, then localizes truth one level up (Tarski hierarchy)
Intuitionistic / type-theoretic	excluded middle; unrestricted Type : Type	bans self-typing — the universe hierarchy; Girard’s paradox is the diagonal of naïve type theory (Coquand [3]; Hurkens [4])
Quantum (orthomodular)	distributivity	the self-negating join lives as a genuine element; complementation stays total
ZTL	identity $p \rightarrow p$, excluded middle, double negation	pointwise quarantine: the liar has no fixed value; it is flagged, not exploded [15]

Two honesty points. First, only the ZTL row is the author’s; the other three are the standard content of their fields, curated here, not discovered. Second, the classical row is the informative one: classical logic retains *every* structural principle, which is exactly why it is the system that must *explode* on the unlocalized liar and repair itself from the metalanguage. Retaining everything is paid for in the hierarchy.

4 Two invariants

Across the four, two things do not move — a dual pair, one positive and one negative.

The floor (a retained principle). Non-contradiction and the reductio it licenses, $(p \rightarrow \neg p) \rightarrow \neg p$, survive in every corner. In the quantum lattice $x \wedge x^\perp = 0$; in the ZTL matrix $\neg(p \wedge \neg p)$ is designated for every value; both are among the propositions of §6.

The ceiling (a retained obstruction). The diagonal survives in every corner, precisely because each logic is expressive enough to state it. What differs is only the *manner of housing* it (§3).

The two are linked: **non-contradiction is the guard against the diagonal**. A logic that surrendered non-contradiction would not tame the diagonal but be trivialized by it (*ex contradictione quodlibet*). Where the obstruction persists, the floor stands; that is why both are invariant.

5 No consistent amalgamation

Why an irreducible plurality rather than a hierarchy or a common refinement?

Because the retained principles are *pairwise incompatible*. Quantum logic retains $x \vee x^\perp = 1$; ZTL withholds $p \vee \neg p$ at the mark — no single valuation structure validates both a total complementation law and its pointwise failure. Type theory demands strong normalization (every term terminates); the ZTL solver houses non-termination as the transient phase N (§6). Any system retaining *all four* families of principles at once would validate distributivity *and* its negation, hence be trivial — that system is classical logic, which pays for retaining everything with the raw diagonal it must then hierarchically contain. There is therefore no consistent common refinement; the four are genuinely incomparable.

We describe this informally as an *orbit* around classical logic — the unreachable centre that keeps all principles at the price of the liar. We stress that “orbit”, “centre” and “cycle” are metaphor: the content is the pairwise incompatibility above and the concrete complementarity of §6. No claim of a formal cyclic structure is made.

6 The ZTL–quantum mirror, machine-checked

The one formalizable part of the picture is the *complementary sacrifice* made by ZTL and by orthomodular logic. Both are recorded below as propositions verified in Lean 4 over finite carriers by `decide`, on the **empty axiom list**, in bare Lean (no `mathlib`, no imports).

Pole A (ZTL). Let $M = \{\mathsf{T}, \mathsf{F}, \mathsf{Z}\}$, with connectives generated by the *zero-trust lift*: for a Boolean f , $\text{lift}(f)(\vec{x}) = \mathsf{T}$ iff f returns **true** under every classical reading of the arguments (Z read as both **true** and **false**), else F . Then in M :

Proposition 6.1 (`lean/ZTL.lean`, empty axiom list).

- (i) non-contradiction — $\neg(p \wedge \neg p) = \mathsf{T}$ for all p ;
- (ii) distributivity in both directions — $p \wedge (q \vee r) = (p \wedge q) \vee (p \wedge r)$ and dually, for all p, q, r ;
- (iii) failure of identity, excluded middle and double negation — $p \rightarrow p$, $p \vee \neg p$, $\neg\neg p \rightarrow p$ are not designated for all p (each is falsified at $p = \mathsf{Z}$).

Here the third symbol Z is a **status mark**, not a truth value: verdicts are two-valued, Z marking an unverified input (greedy register) or, reused positionally in the solver, the transient phase N = “not yet” (Kleene’s undefined), which never labels an input. ZTL is thus **paracomplete** (truth-value gaps, no gluts), not paraconsistent [15].

Pole B (orthomodular). Let MO_2 be the smallest non-distributive orthomodular lattice: $\{0, 1\}$ together with two complementary pairs $(a, a^\perp)$, $(b, b^\perp)$, all four proper elements mutually incomparable, with $x \wedge x^\perp = 0$, $x \vee x^\perp = 1$. Then:

Proposition 6.2 (`lean/QuantumWitness.lean`, empty axiom list).

- (i) non-contradiction — $x \wedge x^\perp = 0$ for all x ;
- (ii) excluded middle and double negation — $x \vee x^\perp = 1$ and $x^{\perp\perp} = x$ for all x ;
- (iii) failure of distributivity — $a \wedge (b \vee b^\perp) = a \wedge 1 = a$, whereas $(a \wedge b) \vee (a \wedge b^\perp) = 0 \vee 0 = 0$.

The mirror. ZTL retains distributivity and surrenders excluded middle, double negation and identity; orthomodular logic retains excluded middle and double negation and surrenders distributivity; both retain non-contradiction. Informally: ZTL fails where a proposition meets *itself*; quantum logic fails where propositions *combine*.

Remark (the honest mark). The two “retentions of excluded middle” are not instances of one theorem in one system. In MO_2 , $x \vee x^\perp = 1$ is a *lattice identity* and $x^{\perp\perp} = x$ holds because the orthocomplement is involutive *by the choice of an ortholattice* — mild, structural. In M , the *failure* of $p \vee \neg p$ is a computation in a *logical matrix*, where validity means designation. These are facts in two different formalisms that rhyme; the “mirror” is an analogy made precise on each side, not a single cross-formalism theorem. This is the boundary between what is proved (each proposition) and what is read (their pairing).

The pair extension. On TWO quanta the mirror sharpens. At the singlet state, pair propositions are true (“spins are opposite”, “total spin = 0”) while *every* local proposition — either particle, any axis — has probability $\frac{1}{2}$: a true join whose joinees are all unverified. The law that falls is the covering law $a \vee b$ true $\iff a$ true or b true — a law the ZTL matrix keeps as a theorem (`cover_or_T`). Lattice expressibility, by contrast, does NOT fall: the singlet line is reachable from local propositions by meet and join (measured in the repository stand `dilemmas/quantum_pair.py`); what is not compositional is the truth VALUE, not the lattice. The seam itself is kernel-checked (`lean/JunctionWitness.lean`, empty axiom list, exact integer arithmetic): the *junction theorem* — the singlet lies in the join of the two product atoms yet in neither atom and in no local plane of either factor. Plugging two quantum logics together makes the COVER face of the disjunction (one face of the trio ZTL surrenders) fall for quantum at the seam of composition, in correlation form — entanglement as a lattice fact; the LEM identity $x \vee x^\perp = 1$ itself stays intact at every size. The falls saturate along the ladder (repository stand `dilemmas/quantum_ladder.py`): one quantum costs distributivity (and modus ponens for the material arrow; the Sasaki arrow keeps both identity and modus ponens), two quanta cost the valuation itself, and beyond two nothing new falls, while $\neg\neg$, non-contradiction and the LEM-identity hold at every size.

Proposition 6.3 (`lean/Contextuality.lean`, empty axiom list).

- (i) Mermin–Peres: *no assignment of ± 1 to the nine cells of the magic square satisfies the six sign constraints (rows +1; columns +1, +1, -1) — 0 of 512, by kernel enumeration;*
- (ii) GHZ: *no assignment of ± 1 to the six single-particle values satisfies $XXX = +1$, $XYX = YXY = YYX = -1$ — 0 of 64.*

The honest split: Lean proves the COMBINATORIAL impossibility for these sign patterns; that the patterns are quantum’s is direct operator arithmetic, verified in the repository stand. Together with the pair measurement: the local truth-value addresses are not merely empty — no bivalent valuation exists to fill them. The quantum “Z” is an ontic vacancy, not an epistemic flag; this is the precise boundary of the kinship between superposition and ZTL’s mark.

Reproduction (Lean 4.29.1, measured 2026-07-18/19):

```
$ lean lean/ZTL.lean           # 13 objects, no axioms,  ~1.1 s
$ lean lean/QuantumWitness.lean #  5 objects, no axioms,  ~0.5 s
$ lean lean/Contextuality.lean  #  3 objects, no axioms,  ~0.4 s
$ lean lean/JunctionWitness.lean #  8 objects, no axioms,  ~0.2 s
```

Every `#print axioms` line reports “*does not depend on any axioms*”, definitions included: no `Classical.choice`, no `Quot.sound`, no `propext`. The intuitionistic and classical corners are, by contrast, *metatheorems* (underivability of excluded middle; failure of strong normalization for Type : Type; “classical as centre”) and are not formalized here by design; the diagonal itself, ranging over four logics that are not objects of one theory, is likewise not a single formal object.

7 The world is not a corner (a reading)

Everything in this section is interpretation, anchored on the measured cells of §6; no formal claim is made beyond Propositions A–C.

The four corners of §3 are DESCRIPTIONS, and it is the descriptions that pay the laws: distributivity, the covering law, the valuation. What Proposition 6.3 exhibits is not a broken inference

inside any lattice — the lattice is consistent — but the absence of a semantic map: the world declines to be a classical model. The discrepancy is located at the INTERFACE between a logic and its world, and it is closed to every repair: no further verification, no better translation, no finer instrument removes it, because there is no valuation to recover (0 of 512). Where a discrepancy between derivation and measurement is normally an incident to be triaged, here it is a theorem.

The same relocation dissolves the diagonal. Lawvere’s construction needs a fixed-point-free map; for every description’s negation this premise holds — the ZTL tables prove $\neg v \neq v$ for all v , and a subspace never equals its orthocomplement, since no nonzero vector is orthogonal to itself — and so every description must TAME the resulting monster, each paying its price: the classical hierarchy of metalanguages, the type-theoretic tower of universes, ZTL’s completeness tax for the quarantine. But at the level of STATES the premise fails: the flip has a fixed point, and it is physical — $X|+\rangle = |+\rangle$, a state prepared and measured on hardware (repository record, IBM Heron, 2026-07-17). The two negations must not be conflated — the lattice orthocomplement remains fixed-point-free while the state-level flip is not — and precisely this non-conflation is the point: the self-negating fixed point, a paradox in every description, EXISTS in the world as a citizen, and its existence costs the world nothing — no hierarchy, no tower, no tax. The paradox is a disease of descriptions, not of the world.

Read this way, the cycle closes cleanly: the world supplies the obstruction (the interface discrepancy, permanent by Proposition 6.3) and natively houses the fixed point; the four logics are four ways of living with that supply, each paying exactly one law for its manner of description. The universe is not a fifth corner. It is what the corners are about.

8 Provenance

	Content	Status
Established	Lawvere’s fixed-point theorem and the diagonal unification (Lawvere [8]; Yanofsky [14]); quantum logic and non-distributivity (Birkhoff–von Neumann [1]); contextuality (Kochen–Specker [7]; Mermin [10]; Peres [13]); intuitionism, the universe hierarchy, Girard’s paradox (Martin-Löf [11]; Coquand [3]; Hurkens [4]); universal logic as a field (Béziau [16])	literature
Author’s	ZTL, the zero-trust paracomplete logic [15]; the arrangement of the four as resolutions of one diagonal; the four Lean witnesses on the empty axiom list (the two poles, the contextuality core, the junction theorem); the measured ladder of falls and its saturation	synthesis + machine-checked
Analogy (no proof)	the “cycle / orbit / centre”; “the floor guards the ceiling”; “each logic pays one principle so none pays the whole diagonal”; §7 entire — “the world is not a corner; the paradox is a disease of descriptions”	structural metaphor, marked

The defensible claim is narrow: four established non-classical logics align as four disciplined res-

olutions of one pre-classical diagonal, and the sharpest pair among them — ZTL and orthomodular logic — make complementary, machine-checkable sacrifices. The organizing picture is offered as a reading; its two poles are offered as proof.

References

- [1] G. Birkhoff, J. von Neumann. *The logic of quantum mechanics*. Ann. of Math. **37** (1936), 823–843.
- [2] D. A. Bochvar. *On a three-valued logical calculus and its application to the analysis of the paradoxes of the classical extended functional calculus*. Mat. Sbornik **4** (1938); Eng. transl. Hist. Philos. Logic **2** (1981), 87–112.
- [3] T. Coquand. *An analysis of Girard’s paradox*. LICS 1986, 227–236.
- [4] A. G. Hurkens. *A simplification of Girard’s paradox*. TLCA 1995, LNCS 902, 266–278.
- [5] S. C. Kleene. *Introduction to Metamathematics*. North-Holland, 1952.
- [6] S. Kripke. *Outline of a theory of truth*. J. Philos. **72** (1975), 690–716.
- [7] S. Kochen, E. P. Specker. *The problem of hidden variables in quantum mechanics*. J. Math. Mech. **17** (1967), 59–87.
- [8] F. W. Lawvere. *Diagonal arguments and cartesian closed categories*. LNM 92, Springer, 1969, 134–145.
- [9] J. Łukasiewicz. *On three-valued logic* (1920). In: Selected Works, North-Holland, 1970.
- [10] N. D. Mermin. *Simple unified form for the major no-hidden-variables theorems*. Phys. Rev. Lett. **65** (1990), 3373–3376.
- [11] P. Martin-Löf. *Intuitionistic Type Theory*. Bibliopolis, 1984.
- [12] L. de Moura, S. Ullrich. *The Lean 4 theorem prover and programming language*. CADE 2021, LNCS 12699, 625–635.
- [13] A. Peres. *Incompatible results of quantum measurements*. Phys. Lett. A **151** (1990), 107–108.
- [14] N. S. Yanofsky. *A universal approach to self-referential paradoxes, incompleteness and fixed points*. Bull. Symbolic Logic **9** (2003), 362–386.
- [15] V. Reznik. *ZTL — Zero-Trust Logic*. Preprint, DOI 10.5281/zenodo.21318981 (concept).
- [16] J.-Y. Béziau (ed.). *Universal Logic: An Anthology*. Birkhäuser, 2012.

AI disclosure. This note was written in dialogue with the AI system Claude (Anthropic); all design decisions, the framing and final responsibility are the human author’s. The reliability of §6 depends on neither the author nor the AI: both Lean files verify against the Lean 4 kernel on the empty axiom list. Everything outside §6 is offered as a reading and marked as such.

License. Text CC BY 4.0; the accompanying Lean files are MIT-licensed (repository ZTL).